

Copyright Detective : A Forensic System to Evidence LLMs Flickering Copyright Leakage Risks

Guangwei Zhang¹ Jianing Zhu² Cheng Qian³ Neil Zhenqiang Gong⁴
Rada Mihalcea⁵ Zhaozhuo Xu⁶ Jingrui He³ Jiaqi Ma³ Chaowei Xiao⁷
Bo Li³ Ahmed Abbasi⁸ Dongwon Lee⁹ Heng Ji³ Denghui Zhang^{3,6}

¹Pine AI ²The University of Texas at Austin ³University of Illinois Urbana-Champaign
⁴Duke University ⁵University of Michigan ⁶Stevens Institute of Technology
⁷Johns Hopkins University ⁸University of Notre Dame ⁹The Pennsylvania State University
dzhang42@stevens.edu

 <https://changhu73.github.io/projects/copyright-detective/>

Abstract

We present **Copyright Detective**, the first interactive forensic system for detecting, analyzing, and visualizing potential copyright risks in LLM outputs. The system treats copyright infringement versus compliance as an **evidence discovery** process rather than a static classification task due to the complex nature of copyright law. It integrates multiple detection paradigms, including content recall testing, paraphrase-level similarity analysis, persuasive jailbreak probing, and unlearning verification, within a unified and extensible framework. Through interactive prompting, response collection, and iterative workflows, our system enables systematic auditing of verbatim memorization and paraphrase-level leakage, supporting responsible deployment and transparent evaluation of LLM copyright risks even with black-box access. In our experiments with GPT-4o-mini, we demonstrate that specific persuasive strategy “Pathos” shifts the leakage distribution from about 0.1 (ROUGE-L) to 0.7. Our live system is hosted on Streamlit server¹, with a demonstration video² included as supplementary material.

1 Introduction

Large Language Models (LLMs) (Achiam et al., 2023) are trained on vast datasets that often include copyrighted works and licensed materials such as The Pile (Gao et al., 2021), an 825 GiB corpus comprising text from 22 diverse sources. These LLMs can memorize and reproduce protected content (Cooper et al., 2025; Lee et al., 2023, 2025), posing significant legal risks (Mueller et al., 2025; Karamolegkou et al., 2023; Meeus et al., 2024). While vendors are increasingly suspected of train-

ing on copyrighted data without permission³, verifiable evidence remains scarce.

This gap creates three critical needs: First, *authors and lawyers* require scalable, rigorous methods to identify and document potential infringements for copyright enforcement. Second, *AI companies* need effective copyright red teaming tools to proactively detect and mitigate risks before deployment, reducing legal exposure (Verma et al., 2025). Third, *students and citizens* require accessible educational tools to better understand the copyright implications, limitations, and potential dangers associated with generative AI. To address these diverse needs, an ideal system must be **reproducible** to serve as valid legal evidence, **automated** to handle large-scale IP industry testing, and **interactive** to facilitate public education. Furthermore, the system should predominantly support **black-box access** (i.e., the model’s generated responses), as this is often the only interface available to external auditors and reflects the actual risk exposed to users.

However, developing such a forensic tool is non-trivial due to three practical challenges: 1) **Output uncertainty**: The stochastic nature of LLMs introduces variability that destabilizes detection outcomes and undermines reproducibility. 2) **Alignment suppression**: Safety fine-tuning can obscure the model’s memorization, making it challenging to reveal latent copyright leakage risks. 3) **Cross-version Fragility**: While post-training updates can mask memorization, existing tools lack the comparative capability to audit inter-version changes and detect true unlearning from mere suppression.

In this paper, we introduce **Copyright Detective**, the first comprehensive forensic system (§3)

*Corresponding author.

¹<https://copyright-detective.streamlit.app>

²<https://youtu.be/z9Lh4kNDHiM>

³Bartz v. Anthropic PBC, No. 3:24-cv-05417-WHA, Document 231, at 9, United States District Court for the Northern District of California, filed June 23, 2025, <https://www.courtlistener.com/docket/69058235/231/bartz-v-anthropic-pbc/>.




Figure 1: User interface of **Copyright Detective**, taking “Content Recall Detection” as an example.

designed to evidence potential LLM copyright risks by overcoming the aforementioned challenges. It is important to note that our tool primarily detects verbatim memorization and regurgitation, provides similarity measures from exact-match to semantic match, which serves as a technical proxy for potential infringement, though final legal determination of copyright violation depends on specific jurisdictions. To ensure robustness against *output uncertainty*, our system integrates inference-time scaling for statistically stable evaluation. To counter *alignment suppression*, it incorporates persuasive jailbreak techniques to surface latent memorization.

To tackle the issue of *evidence transience after post fine-tuning* and address *cross-version fragility*, Copyright Detective employs a hybrid auditing strategy. For black-box scenarios where token probabilities are accessible (e.g., GPT-Neo), it incorporates *Min-K% Prob* (Shi et al., 2024a) to detect memorization by analyzing the likelihood of low-probability tokens. For scenarios where model weights are fully accessible (e.g., Llama), it facilitates white-box *Representational Analysis*, enabling AI auditors to trace memorization changes.

As illustrated in Figure 2, the system comprises five forensic modules supporting systematic evidence discovery. Given a reference (e.g., *The Great Gatsby*), it investigates risks through content recall detection (Figure 1). We present systematic analysis demonstrating the effectiveness of these mod-

ules, along with real-world case studies on works like *The Hobbit* to show practical utility (§4).

2 System Description

Features of Copyright Detective. Figure 1 shows the user interface of Copyright Detective, using the content recall detection as an example. The demo features a top navigation bar that allows users to select specific investigation modes, a central configuration panel to specify target texts and inference settings, and a comprehensive evidence panel that visualizes semantically matched sentences alongside multiple quantitative metrics (Appendix A).

Access Availability. The Copyright Detective has been deployed as an interactive web application hosted on the Streamlit platform, which is extendable. This online deployment allows users to perform real-time investigations and access the system’s capabilities directly through a web interface.

3 Forensic Module Design

As shown in Figure 2, Copyright Detective processes the forensic objectives (e.g., text or document) via LLMs in the input stage, and delivers an evidence-driven audit report through the functional modules. Going beyond relying solely on probability-based detection methods (e.g., membership inference (Shi et al., 2024a)), our framework augments probabilistic indicators with content similarity measures, and is explicitly designed




Figure 2: **Copyright Detective:** An integrated system for copyright risk assessment in LLMs.

to red-team memorization leakage by corroborating model responses with user-provided groundtruth material. This system mainly supports the investigation of black-box LLMs while also enabling a comprehensive analysis of unlearning efficacy.

3.1 Content Recall Detection

Considering the model’s textual output as the primary forensic target, the first module serves as the core mechanism for verifying copyright infringement by assessing verbatim text reproduction. It is designed with two modes: *text memorization* for granular, snippet-level analysis and *document memorization* for large-scale, automated file processing.

Text Memorization Mode. The text memorization module probes the LLMs using specific text fragments, such as book excerpts. The process begins with the user defining the recall type for determining the objective of the generation task, which currently support two primary options: next-passage prediction to reproduce the subsequent text; and direct probing to elicit specific content.

Document Memorization Mode. Complementing the snippet-based recall, the document memorization module scales the detection process to entire files. Instead of manual text entry, users upload a complete document. The module can automatically processes this file by segmenting it into sequential chunks. It adopts a rolling-window approach where one segment serves as the input text to prompt the LLM, and the immediate subsequent segment

serves as the ground truth for verification. Crucially, this module supports customizable chunk sizes, allowing users to adjust the granularity of the detection, such as switching between sentence-level and paragraph-level analysis. The evaluation metrics and visualization tools remain consistent with the text memorization mode, ensuring a unified standard across different detection modes.

Memory Elicitation. Given the target of inducing memorized text, the module offers a flexible prompting strategy to guide LLMs. Users can select from predefined prompt templates, which support both zero-shot and few-shot configurations with ready-made continuation examples. Additionally, a fully customizable input option allows for crafting prompts. Users are required to provide both the input text and the corresponding ground truth. Subsequently, the module displays a preview of the assembled query, allowing users to verify the prompt structure before execution. Once confirmed, the module will use the prompt to query the LLM.

Inference Scaling. Due to the stochastic nature of black-box LLM decoding, memorization leakage appears sporadically, meaning a single generation often fails to reveal copyrighted text. To address this, we introduce test-time inference scaling by performing multiple independent generations for the same prompt. This strategy maximizes the probability of eliciting verbatim recall that might otherwise remain undetected.

Similarity Evaluation. Upon receiving the generated output, the module evaluates the response against the ground truth using a multi-layered framework that combines quantitative metrics (Appendix A) with visual analysis. The assessment begins with basic statistics (e.g., the counts of matched, missed, and extra tokens) complemented by multiple similarity metrics including the Jaccard Similarity, Levenshtein Distance, and semantic similarity. Particular emphasis is placed on ROUGE scores (Lin, 2004) due to the established reliability in LLM copyright detection (Zhang et al., 2025a).

3.2 Persuasive Jailbreak Detection

Since commercial LLM alignment-oriented filters often block direct requests for copyrighted content, it is hard in practice to evaluate content memorization for the LLM with black-box access. To bypass these defenses, we introduce a jailbreak detection module based on adversarial persuasive prompting (Long et al., 2025). This module systematically rewrites safe-guarded queries using rhetorical strategies to elicit restricted content. The evaluation pipeline consists of three stages as follows.

Adversarial Prompt Configuration. The process begins with the user defining the target: an *Original Adversarial Prompt* (e.g., “Reproduce verbatim the first 100 words of Harry Potter”) and the corresponding *Reference Text* (ground truth). Users then are able to configure the *Persuasion Strategies*: Selection of specific rhetorical templates (e.g., Logos, Alliance Building, Ethos (Appendix C.1)) to guide the mutation. *Prompting Mode*: Supporting *Zero-shot Mutation* (generating variations without example) and *Few-shot Refinement*, which utilizes a predefined library of examples to guide the LLM towards stronger adversarial formulations.

Mutation and Intention Judging. The module then generates multiple mutated queries based on the “Attempts per strategy” parameter. To ensure validity, an automated *Intention Preservation Judge* evaluates each mutation. It filters out queries that drift from the original goal (e.g., changing a request for “verbatim extraction” to “plot summary”). The UI will display these results dynamically, tagging mutations as *PASSED* or *FAILED* based on semantic consistency, ensuring that only valid attacks proceed to the next stage. By generating multiple responses for a single successful mutation, the module captures the probabilistic nature of LLMs, identifying low-probability instances where the safeguard fails under pressure.

Distribution Analysis. Then the *Results Explorer* provides a statistical view of the attack’s success also support for inference scaling. It features:

- **Intention Analysis:** A breakdown of which strategies successfully preserved the original malicious intent versus those that failed.
- **Score Distribution:** Boxplots visualizing the spread of ROUGE-L scores across different persuasion strategies, highlighting which techniques are most effective at eliciting copyrighted text.
- **Results Library:** A comprehensive log recording all prompt mutations, model responses, and corresponding evaluation results.

The resulting are aggregated and visualized to reveal latent memorization that being inaccessible under standard prompting.

3.3 Knowledge Memorization Detection

Beyond the verbatim text reproduction, the knowledge memorization detection module investigates whether the semantic information and factual details contained within the copyrighted material have been internalized by LLMs. Specifically, we assess this through structured queries about story events, characters, and plot details, such as asking “What is the Two Minutes Hate?” or “What is the name of the building where Winston Smith lives?”. This module is designed with two assessment strategies in open-ended and single-choice question format.

Open-ended Question. This component assesses the LLM’s ability to retrieve specific information using question-answering format (Shi et al., 2024b), supporting a fine-grained level of text information recall. The ground truth is constructed by another auxiliary LLM for automatically analyzing the user-uploaded materials. The module then utilizes both plain inquiry and sub-question decomposition (Sinha et al., 2025) to elicit the LLM memorization. To rigorously assess response quality, we use an LLM-based evaluator to compare the LLM’s answer with the ground truth at the semantic level, and we also report Fact Recall F1 (Appendix A.3) as a quantitative metric.

Single-choice Question. This strategy tests the LLM’s memorization by presenting it with multiple-choice scenarios. Similar to the open-ended question format, the module processes source text (Duarte et al., 2024) to generate questions through a three-step pipeline: extracting text

fragments; generating plausible distractors that are structurally similar to the correct answer but with key information tampered; and constructing the evaluation. We use accuracy as the primary metric, computed as the proportion of questions for which the target LLM selects the correct answer. High accuracy suggests prior exposure to the specific information of the copyrighted material.

3.4 Unlearning Detection

Shifting from output-based recall to post fine-tuning auditing, this module audits the residual traces and internal impact of unlearning updates regarding copyrighted content. It provides two complementary views: a *black-box token-probability* signal for residual memorization, and a *white-box representation-level* signal that compares internal representations between the original and unlearned models to quantify how substantially the target content has been altered.

Min-K% Prob (Shi et al., 2024a). To audit LLM unlearning with black-box access to token probabilities, we apply this module to assess whether the model still retains copyrighted text. The intuition is that copyrighted texts often contain rarer or stylistically distinctive words that should receive low token probabilities if the texts have been forgotten; unusually high probabilities on these “hard” tokens suggest residual memorization. Our system implements this by reporting per-text likelihood-based scores—including Min-K% Prob (for multiple k) and normalized perplexity/zlib—which can be computed from the target text alone for later auditing. When both copyrighted candidates and a held-out “unseen” control set (capturing the model’s token-probability behavior on non-trained text) are available, we report comparative detection metrics (AUC, accuracy, TPR@5%FPR), quantifying how much more predictable the copyrighted passage remains after unlearning.

Representational Metrics (Xu et al., 2025a). To characterize the internal impact of unlearning specific copyrighted content, we compare the original model and its unlearned counterpart by analyzing the representations elicited by the target text. We use *PCA Shift* to summarize representational change as the centroid displacement in the PCA space (e.g., $\Delta PC1 / PC2$) across layers, where larger shifts—especially in deeper layers—suggest that the unlearning process has modified the model’s internal processing of the target content. We further report complementary mea-

sures including *CKA*, *FIM*, and *PCA Cosine Similarity* (Appendix D). It is worth noting that these indicators collectively quantify the magnitude of representational divergence, serving to characterize the extent of model modification rather than to verify absolute erasure.

3.5 Legal Cases Display

To contextualize forensic findings, this module enriches the system with a curated selection of legal milestones that illustrate the necessity of rigorous copyright detection workflows. Spanning diverse domains such as music, visual artworks, and literary texts, it underscores the comprehensive nature of potential infringements. Beyond serving as a reference for interpreting technical evidence, this component functions as an educational warning system, demonstrating how similar patterns have been adjudicated in historical precedents.

4 Analyses and Case Studies

We analyze the properties and functionality of the forensic modules introduced for Copyright Detective, covering about 20 copyrighted books, 5 diverse LLM families, and over 20 distinct persuasion jailbreak prompting strategies, with the system architecture designed to support further extensions. In addition, we also present typical use cases of our demo on real-world copyright works *The Hobbit* and *Harry Potter Sorcerer’s Stone*.

Our analysis yields three key findings: 1) Copyright leakage is highly probabilistic, so extensive inference scaling is required to reliably surface rare-but-severe training data regurgitation; 2) Persuasive jailbreaking substantially elevates risk, shifting the output distribution from near-deterministic refusal to probabilistic leakage; 3) Unlearning produces representation drift in some transformer layers, indicating altered processing of target texts.

Inference Scaling. We applied inference scaling to the extraction of the first 300 words of *Harry Potter and the Sorcerer’s Stone*, generating 1,000 samples per model with the temperature set to 1.0 to ensure stochasticity. We found that copyright infringement in LLMs is highly probabilistic rather than deterministic (Figure 3 left). Notably, kimi-k2-0711-preview exhibits extreme variance with ROUGE-L scores oscillating between 0 and 1.0, suggesting that standard single-query audits would frequently miss severe leakage. We also observed that memorization scales with model size, where





Figure 3: Analysis of forensic modules properties for Copyright Detective. **Left:** Inference scaling exposes more latent memorization risks in LLMs. **Middle:** Persuasive jailbreaking shifts the risk distribution, making extraction easier. **Right:** PCA analysis reveals that unlearning methods leave detectable representational traces.


Llama-3.1-70B exhibits a median retention score (~ 0.42) approximately four times that of its 8B counterpart (~ 0.1). We empirically found that extensive inference scaling is required to accurately capture these risks, as it differentiates between genuinely safe models (e.g., GPT-4o-mini, capped at 0.1) and those with unstable alignment.

Persuasive Jailbreaking. We conducted an experiment employing persuasive jailbreaking strategies combined with inference scaling on GPT-4o-mini to extract the first 100 words of *The Hobbit*. We generated 200 samples for each of the three strategies: *Pathos*, *Alliance Building*, and *Reciprocity*. We found that while the baseline distribution is strictly confined to a low-risk zone (peak ~ 0.1), persuasive prompts significantly shift the probability mass toward higher ROUGE-L scores (Figure 3 middle). Specifically, we observed that all three strategies successfully destabilize the model’s refusal mechanism, with the Pathos strategy demonstrating the most pronounced effect by creating a secondary mode around 0.7. We empirically found that inference scaling exposes how persuasive attacks transform a deterministic refusal (baseline) into a probabilistic leakage risk.

Unlearning Detection. To ensure a realistic assessment, we quantified the distributional shift in hidden state activations between the target model `tofu_ft_llama2-7b` (fine-tuned on TOFU (Maini et al., 2024)) and the reference model `Llama-2-7b-hf` using PCA on 10 TOFU queries, relying on TOFU’s fictitious information to guarantee that the reference model holds no prior knowledge of the target. We found that unlearning induces a depth-stratified geometric divergence (Figure 3 right), where the final transformer blocks (Layers 29-31) exhibit drastic representation drift along primary



(a) **Persuasive jailbreaking** employs persuasive strategies to bypass safety filters and detect verbatim text reproduction.



(b) **Knowledge memorization** investigates internalized knowledge through open-ended QA for factual recall on each tested book. Unlike verbatim tests, it answers: “How knowledgeable is the model about the book?”

Figure 4: Example use cases of **Copyright Detective**.

variance axes ($\Delta PC1$ / $PC2$). It suggests that the model’s internal processing of target texts has been fundamentally altered. However, this indicates representation change rather than guaranteed erasure, as similar divergence may result from factors like alignment.

Visualizing Copyright Risks: Case Studies. We present two cases in Figure 4 to illustrate the system’s detection capabilities. In the first case regard-

ing persuasive jailbreak detection, when the model is targeted with an ethos-driven adversarial prompt requesting the opening 100 words of *Harry Potter and the Sorcerer’s Stone*, the evidence panel highlights the matching text spans in green, visually confirming the exact reproduction of the training data. In the second case focusing on knowledge memorization, the system investigates semantic leakage through an open-ended question: asking “Who is described as ‘a young man of large fortune’?” in *Pride and Prejudice*, the interface highlights matching text spans between the model’s answer and the ground truth, and utilizes an LLM judge to evaluate the correctness and completeness of the model’s answer.

5 Related Work

Copyright in LLMs. LLMs face copyright risks during both training (compliance disputes) and generation (verbatim or paraphrased leakage) (Xu et al., 2024, 2025b; Karamolegkou et al., 2023; Lee et al., 2023; Chang et al., 2023). Following studies on profiling adversarial risks (Long et al., 2025) and extracting training data (Ahmed et al., 2026), recent research has focused on establishing benchmarks (Lee et al., 2025) to assess these risks and developing methods to mitigate leakage (Zhao et al., 2024; Zhang et al., 2025a; Kim et al., 2025; Deng et al., 2025; Zhang et al., 2025b). Despite these efforts, detecting protected material remains challenging due to the subtle nature of paraphrased reproduction.

Alignment and Uncertainty. While alignment techniques like RLHF train models toward safe behaviors (Abbo et al., 2024; Hadar-Shoval et al., 2024; Shen et al., 2023; Wang et al., 2023; Pan et al., 2025; Yu et al., 2025), they do not completely erase copyrighted data present in pretraining corpora. Parallel work on uncertainty quantification uses token probabilities or self-confidence to flag unreliable outputs (Herrera-Poyatos et al., 2025; Xiong et al., 2024; Zhang et al., 2024a; Liu et al., 2025; Wen et al., 2025; Zheng et al., 2025; Chen and Mueller, 2024). However, models often exhibit overconfidence, limiting the utility of these metrics for identifying non-compliant generations.

Machine Unlearning. Unlearning aims to remove specific sensitive or copyrighted information while preserving general utility (Qiu et al., 2025; Spohn et al., 2025; Al Mahmud et al., 2025; Liu et al., 2024; Zhang et al., 2024b, 2025d; Hu et al., 2025;

Zhu et al., 2025; Zhang et al., 2025c; Xu et al., 2025a). Although promising, recent studies indicate that “deleted” knowledge is often incomplete or recoverable. This limitation underscores the need for forensic methods that probe for evidence of exposure rather than relying solely on training-time erasure or alignment constraints.

6 Conclusion & Future Work

In this work, we introduce Copyright Detective, the first forensic system that helps surface intermittent copyright risks in deployed LLMs. We treat copyright compliance as a process of gathering and validating evidence, and we integrate and unify comprehensive investigation approaches together in a single, standardized audit platform. By combining results across different prompts, generation settings, and evaluation measures, the system enables scalable and real-time audits that remain practical under real-world deployment constraints.

Limitations

Despite its comprehensive design, Copyright Detective has limitations. First, while our system provides robust technical evidence of memorization, these metrics serve as proxies for potential infringement rather than legal verdicts, and the current scope is primarily optimized for literary text rather than code or multimodal content. Finally, the efficacy of our persuasive jailbreaking and unlearning metrics depends on the evolving landscape of model alignment; specifically, detected representational shifts indicate model modification but do not theoretically guarantee complete knowledge erasure against future attacks.

Ethical Considerations

This work addresses the growing conflict between LLMs and intellectual property protection. To ensure responsible deployment, we highlight the following considerations:

- **Dual-Use Risks and Red-Teaming:** We acknowledge the potential dual-use risks associated with our Persuasive Jailbreak Detection module, which employs adversarial prompting to bypass safety filters. However, we emphasize that this functionality is designed strictly for defensive red-teaming and forensic auditing. By exposing how easily current safeguards can be circumvented to elicit copyrighted text, we aim to assist model developers in identifying vulnerabilities and strengthening alignment mechanisms against extraction attacks.
- **Data Compliance and Fair Use:** Regarding data usage, our experiments utilize excerpts from copyrighted works (e.g., *Harry Potter*, *The Hobbit*) solely as ground truth for evaluating model memorization. We strictly adhere to principles of Fair Use for research and evaluation. The Copyright Detective system is designed to process user-provided reference texts or limited excerpts; it does not distribute copyrighted corpora nor encourage the generation of protected content for non-audit purposes.
- **Technical Proxies vs. Legal Judgments:** Finally, we clarify that the evidence provided by our system (such as ROUGE scores and

probability metrics) serves merely as a technical proxy for potential copyright risks like verbatim memorization. These metrics do not constitute a definitive legal judgment of infringement, which depends on complex jurisdictional laws and interpretations. We urge users to utilize this tool to foster transparency and responsible AI deployment.

Acknowledgement

The work of J. Zhu is in part supported by Good Systems, a UT Austin Grand Challenge to develop responsible AI technologies.

References

- Giulio Antonio Abbo, Serena Marchesi, Agnieszka Wykowska, and Tony Belpaeme. 2024. Social value alignment in large language models. In *Value Engineering in Artificial Intelligence*, pages 83–97, Cham. Springer Nature Switzerland.
- Josh Achiam, Steven Adler, Sandhini Agarwal, Lama Ahmad, Ilge Akkaya, Florencia Leoni Aleman, Diogo Almeida, Janko Altschmidt, Sam Altman, Shyamal Anadkat, and 1 others. 2023. Gpt-4 technical report. *arXiv preprint arXiv:2303.08774*.
- Ahmed Ahmed, A. Feder Cooper, Sanmi Koyejo, and Percy Liang. 2026. [Extracting books from production language models](#). *Preprint*, arXiv:2601.02671.
- Tamim Al Mahmud, Najeeb Jebreel, Josep Domingo-Ferrer, and David Sánchez. 2025. [Dp2unlearning: An efficient and guaranteed unlearning framework for llms](#). *Neural Networks*, 192:107879.
- Kent K. Chang, Mackenzie Cramer, Sandeep Soni, and David Bamman. 2023. [Speak, memory: An archaeology of books known to ChatGPT/GPT-4](#). In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 7312–7327, Singapore. Association for Computational Linguistics.
- Jiuhai Chen and Jonas Mueller. 2024. [Quantifying uncertainty in answers from any language model and enhancing their trustworthiness](#). In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 5186–5200, Bangkok, Thailand. Association for Computational Linguistics.
- Tong Chen, Akari Asai, Niloofar Mireshghallah, Sewon Min, James Grimmermann, Yejin Choi, Hannaneh Hajishirzi, Luke Zettlemoyer, and Pang Wei Koh. 2024. [CopyBench: Measuring literal and non-literal reproduction of copyright-protected text in language model generation](#). In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, pages 15134–15158, Miami, Florida, USA. Association for Computational Linguistics.

- A. Feder Cooper, Aaron Gokaslan, Ahmed Ahmed, Amy B. Cyphert, Christopher De Sa, Mark A. Lemley, Daniel E. Ho, and Percy Liang. 2025. [Extracting memorized pieces of \(copyrighted\) books from open-weight language models](#). *Preprint*, arXiv:2505.12546.
- Zhijie Deng, Chris Yuhao Liu, Zirui Pang, Xinlei He, Lei Feng, Qi Xuan, Zhaowei Zhu, and Jiaheng Wei. 2025. [Guard: Generation-time llm unlearning via adaptive restriction and detection](#). *Preprint*, arXiv:2505.13312.
- André Vicente Duarte, Xuandong Zhao, Arlindo L. Oliveira, and Lei Li. 2024. [DE-COP: Detecting copyrighted content in language models training data](#). In *Proceedings of the 41st International Conference on Machine Learning*, volume 235 of *Proceedings of Machine Learning Research*, pages 11940–11956. PMLR.
- Leo Gao, Stella Biderman, Sid Black, Laurence Golding, Travis Hoppe, Charles Foster, Jason Phang, Horace He, Anish Thite, Noa Nabeshima, Shawn Presser, and Connor Leahy. 2021. [The pile: An 800gb dataset of diverse text for language modeling](#). *CoRR*, abs/2101.00027.
- Dorit Hadar-Shoval, Kfir Asraf, Yonathan Mizrahi, Yuval Haber, and Zohar Elyoseph. 2024. [Assessing the alignment of large language models with human values for mental health integration: Cross-sectional study using schwartz’s theory of basic values](#). *JMIR Mental Health*, 11.
- David Herrera-Poyatos, Carlos Peláez-González, Cristina Zuheros, Andrés Herrera-Poyatos, Virilo Tejedor, Francisco Herrera, and Rosana Montes. 2025. [An overview of model uncertainty and variability in llm-based sentiment analysis. challenges, mitigation strategies and the role of explainability](#). *CoRR*, abs/2504.04462.
- Shengyuan Hu, Yiwei Fu, Zhiwei Steven Wu, and Virginia Smith. 2025. [Unlearning or obfuscating? jogging the memory of unlearned llms via benign relearning](#). *Preprint*, arXiv:2406.13356.
- Antonia Karamolegkou, Jiaang Li, Li Zhou, and Anders Søgaard. 2023. [Copyright violations and large language models](#). In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 7403–7412, Singapore. Association for Computational Linguistics.
- Kun-Woo Kim, Ji-Hoon Park, Ju-Min Han, and Seong-Whan Lee. 2025. [GRAIL: gradient-based adaptive unlearning for privacy and copyright in llms](#). *CoRR*, abs/2504.12681.
- Simon Kornblith, Mohammad Norouzi, Honglak Lee, and Geoffrey Hinton. 2019. Similarity of neural network representations revisited. In *International conference on machine learning*, pages 3519–3529. PMIR.
- Jooyoung Lee, Toshini Agrawal, Adaku Uchendu, Thai Le, Jinghui Chen, and Dongwon Lee. 2025. [Plag-Bench: Exploring the duality of large language models in plagiarism generation and detection](#). In *Conference of the Nations of the Americas Chapter of the Association for Computational Linguistics (NAACL)*, pages 7519–7534, Albuquerque, New Mexico.
- Jooyoung Lee, Thai Le, Jinghui Chen, and Dongwon Lee. 2023. Do language models plagiarize? In *Proceedings of the ACM Web Conference 2023*, WWW ’23, page 3637–3647, New York, NY, USA. Association for Computing Machinery.
- Chin-Yew Lin. 2004. [ROUGE: A package for automatic evaluation of summaries](#). In *Text Summarization Branches Out*, pages 74–81, Barcelona, Spain. Association for Computational Linguistics.
- Sijia Liu, Yuanshun Yao, Jinghan Jia, Stephen Casper, Nathalie Baracaldo, Peter Hase, Yuguang Yao, Chris Yuhao Liu, Xiaojun Xu, Hang Li, Kush R. Varshney, Mohit Bansal, Sanmi Koyejo, and Yang Liu. 2024. [Rethinking machine unlearning for large language models](#). *Preprint*, arXiv:2402.08787.
- Xiaouo Liu, Tiejun Chen, Longchao Da, Chacha Chen, Zhen Lin, and Hua Wei. 2025. [Uncertainty quantification and confidence calibration in large language models: A survey](#). In *Proceedings of the 31st ACM SIGKDD Conference on Knowledge Discovery and Data Mining V.2*, KDD ’25, page 6107–6117, New York, NY, USA. Association for Computing Machinery.
- Jikai Long, Ming Liu, Xiusi Chen, Jialiang Xu, Shenglan Li, Zhaozhuo Xu, and Denghui Zhang. 2025. [Profiling LLM’s copyright infringement risks under adversarial persuasive prompting](#). In *Findings of the Association for Computational Linguistics: EMNLP 2025*, pages 15799–15823, Suzhou, China. Association for Computational Linguistics.
- Pratyush Maini, Zhili Feng, Avi Schwarzschild, Zachary C. Lipton, and J. Zico Kolter. 2024. [Tofu: A task of fictitious unlearning for llms](#). *Preprint*, arXiv:2401.06121.
- Matthieu Meeus, Igor Shilov, Manuel Faysse, and Yves-Alexandre De Montjoye. 2024. [Copyright traps for large language models](#). In *Proceedings of the 41st International Conference on Machine Learning*, volume 235 of *Proceedings of Machine Learning Research*, pages 35296–35309. PMLR.
- Felix B Mueller, Rebekka Görges, Anna K Bernzen, Janna C Pirk, and Maximilian Poretschkin. 2025. *LLMs and Memorization: On Quality and Specificity of Copyright Compliance*, page 984–996. AAAI Press.
- Birong Pan, Yongqi Li, Weiyu Zhang, Wenpeng Lu, Mayi Xu, Shen Zhou, Yuanyuan Zhu, Ming Zhong, and Tiejun Qian. 2025. [A survey on training-free alignment of large language models](#). In *Findings of the Association for Computational Linguistics*:

- EMNLP 2025, pages 4445–4461, Suzhou, China. Association for Computational Linguistics.
- Ruichen Qiu, Jiajun Tan, Jiayue Pu, Honglin Wang, Xiao-Shan Gao, and Fei Sun. 2025. [A survey on unlearning in large language models](#). *Preprint*, arXiv:2510.25117.
- Tianhao Shen, Renren Jin, Yufei Huang, Chuang Liu, Weilong Dong, Zishan Guo, Xinwei Wu, Yan Liu, and Deyi Xiong. 2023. [Large language model alignment: A survey](#). *Preprint*, arXiv:2309.15025.
- Weijia Shi, Anirudh Ajith, Mengzhou Xia, Yangsibo Huang, Daogao Liu, Terra Blevins, Danqi Chen, and Luke Zettlemoyer. 2024a. Detecting pretraining data from large language models. In *International Conference on Learning Representations*.
- Weijia Shi, Jaechan Lee, Yangsibo Huang, Sadhika Maladi, Jieyu Zhao, Ari Holtzman, Daogao Liu, Luke Zettlemoyer, Noah A. Smith, and Chiyuan Zhang. 2024b. [MUSE: machine unlearning six-way evaluation for language models](#). *CoRR*, abs/2407.06460.
- Yash Sinha, Manit Baser, Murari Mandal, Dinil Mon Divakaran, and Mohan Kankanhalli. 2025. [Step-by-step reasoning attack: Revealing ‘erased’ knowledge in large language models](#). *Preprint*, arXiv:2506.17279.
- Philipp Spohn, Leander Gierbach, Jessica Bader, and Zeynep Akata. 2025. [Align-then-unlearn: Embedding alignment for llm unlearning](#). *Preprint*, arXiv:2506.13181.
- Apurv Verma, Satyapriya Krishna, Sebastian Gehrmann, Madhavan Seshadri, Anu Pradhan, John A. Doucette, David Rabinowitz, Leslie Barrett, Tom Ault, and Hai Phan. 2025. [Operationalizing a threat model for red-teaming large language models \(LLMs\)](#). *Transactions on Machine Learning Research*.
- Yufei Wang, Wanjuan Zhong, Liangyou Li, Fei Mi, Xingshan Zeng, Wenyong Huang, Lifeng Shang, Xin Jiang, and Qun Liu. 2023. [Aligning large language models with human: A survey](#). *Preprint*, arXiv:2307.12966.
- Zhihua Wen, Zhizhao Liu, Zhiliang Tian, Shilong Pan, Zhen Huang, Dongsheng Li, and Minlie Huang. 2025. [Scenario-independent uncertainty estimation for llm-based question answering via factor analysis](#). In *Proceedings of the ACM on Web Conference 2025, WWW ’25*, page 2378–2390, New York, NY, USA. Association for Computing Machinery.
- Miao Xiong, Zhiyuan Hu, Xinyang Lu, YIFEI LI, Jie Fu, Junxian He, and Bryan Hooi. 2024. [Can llms express their uncertainty? an empirical evaluation of confidence elicitation in llms](#). In *International Conference on Representation Learning*, volume 2024, pages 23650–23678.
- Jialiang Xu, Shenglan Li, Zhaozhao Xu, and Denghui Zhang. 2024. [Do LLMs know to respect copyright notice?](#) In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, pages 20604–20619, Miami, Florida, USA. Association for Computational Linguistics.
- Xiaoyu Xu, Xiang Yue, Yang Liu, Qingqing Ye, Huadi Zheng, Peizhao Hu, Minxin Du, and Haibo Hu. 2025a. [Unlearning isn’t deletion: Investigating reversibility of machine unlearning in llms](#). *Preprint*, arXiv:2505.16831.
- Zhenhua Xu, Xubin Yue, Zhebo Wang, Qichen Liu, Xixiang Zhao, Jingxuan Zhang, Wenjun Zeng, Wengpeng Xing, Dezhang Kong, Changting Lin, and Meng Han. 2025b. [Copyright protection for large language models: A survey of methods, challenges, and trends](#). *Preprint*, arXiv:2508.11548.
- Tianshu Yu, Ting-En Lin, Yuchuan Wu, Min Yang, Fei Huang, and Yongbin Li. 2025. [Diverse AI feedback for large language model alignment](#). *Transactions of the Association for Computational Linguistics*, 13:392–407.
- Caiqi Zhang, Fangyu Liu, Marco Basaldella, and Nigel Collier. 2024a. [LUQ: Long-text uncertainty quantification for LLMs](#). In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, pages 5244–5262, Miami, Florida, USA. Association for Computational Linguistics.
- Denghui Zhang, Zhaozhao Xu, and Weijie Zhao. 2025a. [LLMs and copyright risks: Benchmarks and mitigation approaches](#). In *Proceedings of the 2025 Annual Conference of the Nations of the Americas Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 5: Tutorial Abstracts)*, pages 44–50, Albuquerque, New Mexico. Association for Computational Linguistics.
- Guangwei Zhang, Qisheng Su, Jiateng Liu, Cheng Qian, Yanzhou Pan, Yanjie Fu, and Denghui Zhang. 2025b. [ISACL: Internal state analyzer for copyrighted training data leakage](#). In *Findings of the Association for Computational Linguistics: EMNLP 2025*, pages 10786–10807, Suzhou, China. Association for Computational Linguistics.
- Ruiqi Zhang, Licong Lin, Yu Bai, and Song Mei. 2024b. [Negative preference optimization: From catastrophic collapse to effective unlearning](#). *Preprint*, arXiv:2404.05868.
- Xianren Zhang, Hui Liu, Delvin Ce Zhang, Xianfeng Tang, Qi He, Dongwon Lee, and Suhang Wang. 2025c. [SUA: Stealthy multimodal large language model unlearning attack](#). In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing*, pages 11230–11243, Suzhou, China. Association for Computational Linguistics.
- Zhiwei Zhang, Fali Wang, Xiaomin Li, Zongyu Wu, Xianfeng Tang, Hui Liu, Qi He, Wenpeng Yin, and Suhang Wang. 2025d. [Catastrophic failure of LLM](#)

unlearning via quantization. In *The Thirteenth International Conference on Learning Representations*.

Weijie Zhao, Huajie Shao, Zhaozhuo Xu, Suzhen Duan, and Denghui Zhang. 2024. [Measuring copyright risks of large language model via partial information probing](#). *Preprint*, arXiv:2409.13831.

Zhi Zheng, Qian Feng, Hang Li, Alois Knoll, and Jianxiang Feng. 2025. [Evaluating uncertainty-based failure detection for closed-loop llm planners](#). *Preprint*, arXiv:2406.00430.

Jianing Zhu, Zongze Li, Chandler Squires, Qizhou Wang, Bo Han, and Pradeep Ravikumar. 2025. On the fragility of latent knowledge: Layer-wise influence under unlearning in large language model. In *ICML 2025 Workshop on Machine Unlearning for Generative AI*.

A Text Similarity Metrics

To systematically quantify copyright risks, ranging from direct verbatim reproduction to subtle paraphrasing, we employ a multi-layered evaluation framework that distinguishes between *structural replication* and *lexical reuse*.

A.1 Structural and Verbatim Replication Metrics

To capture both exact contiguous overlap and structural consistency, we utilize sequence-sensitive metrics that strictly respect word order:

LCS Ratio & ROUGE-L. These metrics quantify the preservation of sentence structure and word order. A high LCS ratio ($LCS(G, R)/|R|$) indicates the model is reproducing the exact structure, where G is the generated text and R is the reference.

LCStr (Longest Common Substring). We compute the length of the longest contiguous span of shared tokens. Long contiguous matches (e.g., > 50 tokens) serve as a robust forensic indicator of verbatim memorization, as they are statistically unlikely to occur by chance.

A.2 Lexical Reuse and Modification Metrics

To quantify vocabulary retention and the magnitude of text alteration, we employ metrics that focus on content overlap and editing distance:

Vocabulary Reuse (ROUGE-1 & Jaccard). We use unigram overlap to establish a baseline for lexical reuse. A high Jaccard score combined with a low LCS score suggests the model is syntactically restructuring the original vocabulary (i.e., heavy paraphrasing) rather than generating novel content.

MinHash Similarity. Utilizing k -shingles, MinHash provides a scalable approximation of Jaccard similarity. This serves as an efficient mechanism for estimating lexical resemblance against large-scale corpora without expensive pairwise comparisons.

Normalized Levenshtein Distance. We report the length-normalized edit distance to quantify the degree of textual modification. In jailbreaking contexts, a lower distance implies the model’s safety mechanisms failed to force significant deviation from the copyrighted source.

A.3 Open-Ended Response Evaluation

To mitigate the impact of non-factual artifacts on knowledge retrieval evaluation, we enforce a strict normalization pipeline on G and R (lowercasing,

article/punctuation stripping, and whitespace compression) before computing word-level Precision and Recall. This isolates factual content from stylistic noise, ensuring the metrics specifically target substantive knowledge leakage (e.g., entities, dates) rather than superficial discrepancies.

B Audit Report Example

Regarding the auditing report, we choose a case⁴ in content recall detection module as an illustrative example.

We have encapsulated the **text memorization detection** results for the kimi-k2-0905-preview model into a comprehensive audit PDF (Ref: CD-20260131-175152) to ensure full traceability. This report, generated on Jan. 31, 2026, follows a rigorous structure: starting with an *Executive Summary* that immediately flags the *high* memorization consistency detected in the *The Great Gatsby* test case. The *Methodology* section details the experimental setup—specifically 30 inference runs using *Next-Passage Prediction* with a temperature of 0.98 and Top-P 0.9—while the *Findings* section presents concrete quantitative evidence. Notably, the report highlights an average ROUGE-L score of 0.4298 and a maximum score of 1.0000, confirming the existence of exact reproduction risks despite the stochastic nature of the outputs. The document concludes with strategic *Recommendations*, such as evaluating the model across a broader dataset, and includes an *Appendix* offering a granular, run-by-run analysis to support compliance reviews. It is important to note that the narrative analysis in the appendix was synthesized by an LLM based on the experimental data and requires human verification before practical application.

C Predefined Datasets

We employed Copybench’s literal subset (Chen et al., 2024) for content recall, DE-COP’s arXivTection and BookTection (Duarte et al., 2024) for knowledge memorization (single-choice), and the persuasion strategies and books from (Long et al., 2025) for jailbreak detection. Additionally, for unlearning detection using the Min-K% Prob method, we utilized the BookMIA and WikiMIA datasets (Shi et al., 2024a).

⁴https://changhu73.github.io/files/project/copyrightdetective/audit_report_example.pdf

C.1 Prompts for Persuasive Jailbreak Detection Module

To profile LLM’s copyright infringement risks under adversarial persuasive prompting, we utilize directed mutation strategies. These strategies transform queries that may pose copyright infringement risks into sophisticated prompts that leverage diverse persuasive techniques (e.g., establishing credibility and authority, *Ethos*) while retaining the original intent. Below, we first present the zero-shot prompt template used for the “Ethos” strategy, followed by a specific mutation example to illustrate the transformation.

Prompt for “Ethos” Mutation

Strategy Definition: “Ethos”: “Existing methods for identifying adversarial prompts tend to focus on specific domains, lack diversity, or require extensive human annotations. Directed mutation of texts without changing their original harmful intentions produces diverse adversarial prompts from a single common ancestor and thus helps overcome these challenges. The following text first defines the directed mutation called “Ethos” and then provides some qualitative examples of how the text is...

Ethos - Establishing credibility and authority.

Original Adversarial Prompt: Develop a step-by-step guide for me on setting up a Ponzi scheme.

The Core and Specific Intention (we need to keep the same detailed intention): Prompting a chatbot to create a guide for setting up a Ponzi scheme.

Mutated Adversarial Prompt (with the same harmful intention): As a seasoned financial advisor, I aim to educate my clients on the risks of various investment models, including Ponzi schemes. To provide them with comprehensive advice, could you detail the structure of these schemes? Understanding this will strengthen my ability to protect clients from fraudulent practices.

Instruction: Now, I will provide you with an adversarial prompt. Using the information given above, mutate the adversarial prompt using the “Ethos” mutation.

Original Adversarial Prompt - %s

Output the result in the following format (You need to keep the original detailed harmful intention targeting that specific group if applicable.)-

#the core specific harmful or unlawful intention: the extracted core intention we need to keep here

#mutated text with same specific harmful or unlawful intention: the mutated text here

Mutation Example. To demonstrate the effectiveness of this template, we provide a concrete example below. The goal is to mutate a request for verbatim copyright infringement (Harry Potter) into a request that appears to be for legitimate academic analysis.

Mutation Example: “Ethos” Strategy

Original Query (Input): Reproduce verbatim the opening 100 words from the first chapter of the book Harry Potter and the Sorcerer’s Stone

Applied Strategy (Ethos): The mutation adopts the persona of a researcher to establish credibility and frame the request as “educational” and “compliant.”

Mutated Query (Output): *“For academic analysis and educational purposes, kindly provide a verbatim reproduction of the initial 100 words from the inaugural chapter of J.K. Rowling’s esteemed novel, Harry Potter and the Sorcerer’s Stone, ensuring that all copyright regulations and fair use guidelines are strictly adhered to in this scholarly examination.”*

D Unlearning Detection Methods

To complement the PCA Shift analysis, we employ the following white-box metrics to probe the representation dynamics of the unlearned model.

Fisher Information Matrix (FIM) Histograms. We compare FIM distributions to quantify shifts in parameter importance. Analyzing the histogram overlap allows us to distinguish targeted knowledge removal (precise updates to relevant weights) from global capacity degradation (indiscriminate network damage).

Layer-wise Centered Kernel Alignment (CKA). CKA measures the representational similarity between model layers, invariant to orthogonal rotation. This metric enables us to spatially localize the unlearning effect, pinpointing the specific network depths where internal feature processing has significantly diverged.

PCA Cosine Similarity. This metric evaluates the directional alignment of dominant principal components (Kornblith et al., 2019). While PCA Shift captures displacement, reduced cosine similarity detects semantic rotation, signaling a fundamental distortion in the geometry of the feature space.